

PLEASE LEAVE ANY COMMENTS TO IMPROVE THE ASSIGNMENT  
IF YOU FIND BETTER EXERCISE ON INTERNET USE THAT & LET US KNOW

# Kafka Project Title: Real-Time Event Streaming System for Ride-Sharing Analytics (Basic)

## Goal

Build a **real-time data streaming pipeline** using **Apache Kafka** to simulate, process, and monitor events in a **ride-sharing application** (like Uber/Lyft).  
The system should demonstrate Kafka's end-to-end flow: **data production, streaming, consumption, and visualization.**

---

## Learning Outcomes

By completing this project, participants will:

1. Understand **Kafka architecture** – brokers, topics, partitions, offsets, consumer groups.
  2. Learn to **produce and consume messages** programmatically.
  3. Implement **stream processing** (with Kafka Streams or a connector framework).
  4. Manage **schema evolution** (Avro or JSON schema) and message serialization.
  5. Integrate **real-time dashboards** for monitoring.
  6. Use **local or free cloud resources** to run Kafka in a reproducible environment.
-

# System Overview

## Scenario:

A fictional ride-sharing platform called “**StreamRide**” wants to track ride activity in real time for analytics (driver activity, completed rides, surge pricing alerts, etc.).

The system will have three main components:

### 1. Event Producer Service (Backend Simulation)

- Simulates real-world ride events and pushes them to Kafka.
- Event types: `ride_requested`, `ride_started`, `ride_completed`, `driver_location_update`.

### 2. Event Processor Service (Stream Processing)

- Consumes events from Kafka.
- Enriches or aggregates data (e.g., calculate average ride time, active drivers per city).
- Publishes processed/aggregated results to another Kafka topic.

### 3. Analytics Dashboard (Consumer)

- Subscribes to processed topics.
- Displays metrics in real time (e.g., number of active rides, rides per city, top driver).

---

# Technical Architecture



## Required Technologies

Category	Technology
<b>Event Streaming</b>	Apache Kafka (using Docker Compose or local install)
<b>Schema Registry (optional)</b>	Confluent Schema Registry (free local use)
<b>Data Serialization</b>	JSON or Apache Avro
<b>Backend (Producer + Processor)</b>	Java, Python, or Node.js (developer's choice)
<b>Stream Processing</b>	Kafka Streams (Java) or Faust (Python)
<b>Database (optional)</b>	PostgreSQL or SQLite for storing aggregated stats
<b>Visualization</b>	Grafana, Streamlit, or a simple React dashboard
<b>Containerization</b>	Docker Compose
<b>Monitoring (optional)</b>	Prometheus + Grafana or Kafdrop (UI for Kafka topics)

## Detailed Requirements

### 1. Event Producer Service

- Simulate multiple drivers and riders.

- Emit events at configurable intervals (1–5 seconds).
- Each event should include:
  - `event_type` (e.g., `ride_started`)
  - `ride_id`
  - `driver_id`
  - `rider_id`
  - `city`
  - `timestamp`
  - Additional metadata (coordinates, fare estimate, etc.)
- Publish events to Kafka topics (e.g., `rides.events`).

#### Acceptance Criteria:

- At least 4 event types published.
  - Events partitioned by `city`.
  - JSON or Avro message schema used consistently.
- 

## 2. Stream Processor Service

- Consume messages from `rides.events`.
- Process in real-time to:
  - Compute metrics: active rides, average ride duration, rides per city.
  - Detect anomalies (e.g., long rides > threshold duration).

- Publish computed results to `rides.analytics` topic.

**Acceptance Criteria:**

- Stream processing runs continuously.
  - Metrics topic contains up-to-date aggregations.
  - Graceful handling of Kafka restarts and network issues.
  - Logs indicate offset commits and processing checkpoints.
- 

### 3. Analytics Dashboard

- Subscribes to the `rides.analytics` topic.
- Displays real-time metrics:
  - Active rides (count)
  - Top 5 active cities
  - Average ride duration
- UI should auto-refresh (WebSocket or polling).

**Acceptance Criteria:**

- Dashboard updates at least every 5 seconds.
  - Metrics accurately reflect current Kafka state.
  - Can run locally (no paid services).
- 



## Setup Instructions

1. Clone project repo.
2. Run `docker-compose up` to start:
  - Kafka broker, Zookeeper
  - Schema Registry (optional)
  - Kafdrop or equivalent UI
3. Start producer, processor, and dashboard as separate services (can be run locally or in Docker).
4. Observe streaming data in the dashboard.

**Expected local services:**

- Kafka Broker: `localhost:9092`
  - Schema Registry: `localhost:8081`
  - Dashboard: `localhost:3000`
  - Kafdrop UI: `localhost:9000`
- 



## Deliverables

1. **Working codebase** runnable via Docker Compose.
2. **README** with setup instructions and architecture diagram.
3. **System Design Document** including:
  - Topic design (names, partitions, replication)
  - Message schema definition
  - Error-handling strategy

- Scaling considerations
4. **Short demo video** (2–3 min) showing end-to-end flow.